


# CS 561: Artificial Intelligence

---

*Local Search*

# Local Search

- Previously: systematic exploration of search space
  - Path to goal is solution to problem
- YET, for some problems path is irrelevant
  - E.g. 8-queens (only final configuration matters and not the path)
- Other applications
  - Integrated circuit design
  - Factory-floor layout
  - Job-shop scheduling
  - Automatic programming
  - Network optimizations
  - Vehicle routing
  - Portfolio management



# Local Search

---

- Different algorithms can be used: **Local Search**
  - Hill-climbing or Gradient ascent
  - Simulated Annealing
  - Genetic Algorithms, others...
- Useful for solving pure optimization problems
  - Systematic search doesn't work
  - Aim is to find best state according to an *objective function*
  - However, can start with a suboptimal solution and improve it
- Many optimization problems don't fit the *standard search technique*
  - Nature provides an *objective function-reproductive fitness* that *Darwinian* evolution could be seen as attempting to *optimize*, but there is no *goal test* and *path cost* for the problem

# Local search and optimization

---

- Local search: *use single current state and move to neighboring states*
- Advantages:
  - Uses very little memory
  - Finds often reasonable solutions in large or infinite state spaces
- Are also useful for pure optimization problems
  - Find best state according to some *objective function*


# Local Search and Optimization

---

- State space landscape
  - Location (defined by state)
- Elevation
  - Defined by the value of the *heuristic function* or *objective function*
- Elevation corresponds to *cost*
  - *Global minimum* (aim is to find the lowest valley)
- Elevation corresponds to an *objective function*
  - *Global maximum* (aim is to find the highest peak)
- Complete algorithm always finds a goal if one exists
- Optimal algorithm always finds global minimum/maximum

# Local search and optimization

---



# Hill-climbing search

---

- “is a loop that continuously moves in the direction of increasing value”
  - It terminates when a peak is reached
  - No neighbor has higher value
- does not maintain any search tree
  - Current node data structure records *state* and *objective function value*
- does not look ahead beyond the immediate neighbors of the current state
- chooses randomly among the set of best successors, if there is more than one

# Hill-climbing search

---

- Hill-climbing a.k.a. **greedy local search**
  - Grabs a good neighbor state without thinking ahead about where to go next
- Makes very rapid progress towards a solution
  - Quite easy to improve a bad state
- Some problem spaces are great for hill climbing and others are terrible



# Hill-climbing search (*steepest ascent*)

---

**function** HILL-CLIMBING(*problem*) **return** a state that is a (global) maximum

**input:** *problem*, a problem

**local variables:** *current*, a node.

*neighbor*, a node.

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])


**loop do**

*neighbor*  $\leftarrow$  a highest valued successor of *current*

**if** VALUE [*neighbor*]  $\leq$  VALUE[*current*] **then return** STATE[*current*]

*current*  $\leftarrow$  *neighbor*


# Hill climbing example



$$f(n) = -(\text{number of tiles out of place})$$

# Local search and optimization

---



# Drawbacks of hill climbing

---

- **Local Maxima**
  - peaks that aren't the highest point in the space (*below global maxima*)
  - higher than each of its neighboring states
- **Plateau:** region where evaluation function is flat
  - **Flat local maximum:** no uphill exit exists
  - **Shoulder:** possible to make progress
  - Could give the search algorithm no direction (*random walk*) for local maximum

# Drawbacks of hill climbing

---

- **Ridges:**

- flat like a plateau, but with drop-offs to the sides; steps to the North, East, South and West may go down, but a combination of two steps (e.g. N, W) may go up
- *results in a sequence of local maximum not connected to each other*

- **For 8-queens problem**

- Hill climbing gets stuck for the 86% problem instances
- Solves only 14% instances
- Average # steps for successful moves: 4
- Average # steps for unsuccessful moves: 3
- acceptable solution for a 17 million states


# Hill Climbing

---

- Algorithm halts
  - Reaches a plateau where the best successor has the same value as the current state
- Solution: *sideways move* in the hope that plateau is really a shoulder
  - Infinite loop could occur whenever the algorithm reaches a flat local maxima that is not shoulder
  - **Solution**: put a limit on the number of consecutive sideways moves allowed
- For the 8-queens problem
  - 100 sideways moves raises the solution level by 80%
  - But with additional cost of roughly 21 steps for each successful instance and 64 for a failure
- Remedy: *Introduce randomness to cope up with the problems*

# Local search and optimization

---



# Hill-climbing variations

---

- Stochastic hill-climbing

- Random selection among the uphill moves
- Selection probability can vary with the steepness of the uphill move
- Converges slowly than steepest ascent (*best first search*) but finds better solution in some cases

- First-choice hill-climbing

- Stochastic hill climbing by generating successors randomly until a better one is found
- Good strategy when a state has many successors

- Random-restart hill-climbing

- Tries to avoid getting stuck in local maxima.
- If at first you don't succeed, try, try again...



# Random-restart hill climbing (*some more details*)

---

- Conducts a series of hill-climbing searches from randomly generated initial states and stop when goal is found

*If there are few local maxima and plateaus, it will find solutions very quickly! (for three million queens solutions can be obtained within a minute)*

# Random hill climbing (*An Example*)

---

- Conducts a series of hill-climbing searches from randomly generated initial states and stop when goal is found

Suppose,  $p$ : probability of success

# restarts =  $1/p$

For 8-queens with no sideways,  $p=0.14$

#iterations required = 7 for finding a goal (failure: 6 and success: 1)

# expected steps = *cost of one successful iteration* +  $(1-p)/p$  \* *cost of failure*

=  $4 + (1-0.14)/0.14 * 3 = 4+18=22$  [4: average # of steps for a successful move; 3: average # of steps for an unsuccessful move]

- Success of hill climbing depends on the shape of the state-space landscape

*If there are few local maxima and plateaus, it will find solutions very quickly!*

# Random hill climbing (*An Example*)

---

- With side-ways moves

$1/0.94=1.06$  iterations required

Steps:

$$1*21 + (0.06/0.94) * 64 = 25$$

Assuming

21: steps for success

64: Steps for failure

---

## ■ Why is hill climbing useful?

- when the amount of time available to perform a search is limited, such as with real-time systems
- can return a valid solution even if it's interrupted at any time before it ends

# Simulated annealing

---

- Hill climbing that does *not make downhill move* is *incomplete*
- Pure random walk: choosing successor from a list is *complete* but *inefficient*

*Solution: hill climbing + random walk = simulated annealing*

*to ensure both efficiency and completeness*

# Simulated Annealing

---

Use a more complex Evaluation Function:

- Do sometimes accept candidates with higher cost to escape from local optimum
  - Idea: but gradually decrease their size and frequency
- Adapt the parameters of this evaluation function during execution
- Based upon the analogy with the simulation of the annealing of solids

# Simulated annealing

---

- Origin: metallurgical annealing
  - Temper or harden metals and glass by heating them to a high temperature and then gradually cooling them
  - Allowing material to coalesce into a low-energy crystalline state
- Application: VLSI layout, airline scheduling, TSP etc.

# Other Names

---

- Monte Carlo Annealing
- Statistical Cooling
- Probabilistic Hill Climbing
- Stochastic Relaxation
- Probabilistic Exchange Algorithm



# Analogy

---

- **Slowly** cool down a heated solid, so that all particles arrange in the ground energy state
- At each temperature wait until the solid reaches its thermal equilibrium
- Probability of being in a state with energy  $E$ :

$$Pr \{ \mathbf{E} = E \} = 1/Z(T) \cdot \exp (-E / k_B \cdot T)$$

$E$       Energy

$T$       Temperature

$k_B$       Boltzmann constant

$Z(T)$       Normalization factor (temperature dependent)

# Simulation of cooling (Metropolis 1953)

---

At a fixed temperature  $T$ :

- Perturb (randomly) the current state to a new state
- $\Delta E$  is the difference in energy between new and current state
- If  $\Delta E < 0$  (new state is lower), accept new state as current state
- If  $\Delta E \geq 0$ , accept new state with probability
$$Pr(\text{accepted}) = \exp(-\Delta E / k_B \cdot T)$$
- Eventually the system evolves into thermal equilibrium at temperature  $T$
- When equilibrium is reached, temperature  $T$  can be lowered and the process can be repeated

# Simulated Annealing

---

- Same algorithm can be used for combinatorial optimization problems:

Energy  $E$  corresponds to the Cost function  $C$

Temperature  $T$  corresponds to control parameter  $c$

$$Pr \{ \text{configuration} = i \} = 1/Q(c) \cdot \exp (-C(i) / c)$$

$C$  Cost

$c$  Control parameter

$Q(c)$  Normalization factor (not important)

# Simulated annealing

---

- Unlike hill climbing, it does not always pick the *best move*
- SA picks the move randomly
  - If situation improves then accept the move
  - Otherwise, accept the move with some probability
- Probability decreases as the temperature goes down
- Probability decreases exponentially with the *badness of the move*
- Bad moves are more likely to be allowed at the start when temperature is high, and they are unlikely when temperature decreases

# Simulated annealing

---

**function** SIMULATED-ANNEALING( *problem*, *schedule*) **return** a solution state

**input:** *problem*, a problem

*schedule*, a mapping from time to temperature

**local variables:** *current*, a node.

*next*, a node.

*T*, a “temperature” controlling the probability of downward steps

*current*  $\leftarrow$  MAKE-NODE(INITIAL-STATE[*problem*])

**for** *t*  $\leftarrow$  1 **to**  $\infty$  **do**

*T*  $\leftarrow$  *schedule*[*t*]

**if** *T* = 0 **then return** *current*

*next*  $\leftarrow$  a randomly selected successor of *current*

$\Delta E \leftarrow$  VALUE[*current*] - VALUE[*next*]

**if**  $\Delta E > 0$  **then** *current*  $\leftarrow$  *next*

**else** *current*  $\leftarrow$  *next* only with probability  $e^{-\Delta E / T}$

# Local beam search

---

- Keep track of  $k$  states instead of one
  - Initially:  $k$  random states
  - Next: determine all successors of  $k$  states
  - If any of successors is goal  $\rightarrow$  finished
  - Else select  $k$  best from successors and repeat.
- *Is it equivalent of running  $k$  random restarts in parallel ?*
- Random-restart search vs. Beam search
  - In random-restart, each process runs independently
  - In beam-search, information is shared among  $k$  search threads

# Local beam search

---

- Kth state generates good successors and all  $k-1$  states generate bad states
  - Abandons  $k-1$  unfruitful searches and move to the  $k$ th state
- Can suffer from lack of diversity
  - Quickly become concentrated in a small region of the state space
  - Not very good in comparison to expensive version of hill climbing
- Stochastic variant: choose  $k$  successors (*at random and not the best ones*) at proportionally to state success
- Resemblance to the process of natural selection
  - *Successors (offspring)*
  - *State (organism)*
  - *Populate the next generation according to fitness value*

# Genetic Algorithm: Introduction

---

- Randomized search and optimization technique
- Evolution produces good individuals, similar principles might work for solving complex problems
- Developed: USA in the 1970' s by J. Holland
- Got popular in the late 1980's
- Early names: J. Holland, K. DeJong, D. Goldberg
- Based on ideas from *Darwinian Evolution*
- Can be used to solve a variety of problems that are not easy to solve using other techniques



# Genetic algorithms

---

- Variant of stochastic beam search
- *Guided by the principles of natural genetics that follow Darwin evolution*
- A successor state is generated by combining two parent states rather than modifying a single state
- Start with  $k$  randomly generated states (**population**)
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s): *chromosome representation*
- Evaluation function (**fitness function**): Higher values for better states
- Produce the next generation of states by *selection*, *crossover*, and *mutation*

# Evolution in the real world

---

- Each cell of a living being contains **chromosomes**-strings of *DNA*
- Each chromosome contains a set of **genes**-blocks of DNA
- Each gene determines some aspects of the organism (like *eye colour*)
- A collection of genes sometimes called a **genotype**
- A collection of aspects (like *eye colour*) sometimes called a **phenotype**
- Reproduction involves recombination of genes from parents and then small amount of **mutation** (errors) in copying
- The **fitness** of an organism is how much it can reproduce before it dies
- Evolution based on “**survival of the fittest**”

# Genetic Algorithm: Similarity with Nature

---

Genetic Algorithms	↔	Nature
A solution (phenotype)		Individual
Representation of a solution ( <i>genotype</i> )		Chromosome
Components of the solution		Genes
Set of solutions		Population
Survival of the fittest ( <i>Selection</i> )		Darwins theory
Search operators		Crossover and mutation
Iterative procedure		Generations

# Genetic Algorithm: Basic Concepts

---

- Parameters of the search space encoded in the form of strings - *Chromosomes*
- Collection of chromosomes - *Population*
- *Initial step*: a random population representing different points in the search space
- *Objective or fitness function*: associated with each string
  - represents the degree of *goodness* of the string

# Genetic Algorithm: Basic Concepts

---

- Selection
  - Based on the principle of survival of the fittest, a few of the strings selected
- Biologically inspired operators like *crossover* and *mutation* applied on these strings to yield a new generation of strings
- Processes of *selection*, *crossover* and *mutation* continue for a fixed number of generations or till a termination condition satisfied

# Basic Steps of Genetic Algorithm

---

1.  $t = 0$
  2. initialize population  $P(t)$  /\*  $Popsiz = |P|$  \*/
  3. for  $i = 1$  to  $Popsiz$   
    compute fitness  $P(t)$
  4.  $t = t + 1$
  5. if termination criterion achieved go to step 10
  6. select ( $P$ )
  7. crossover ( $P$ )
  8. mutate ( $P$ )
  9. go to step 3
  10. output best chromosome and stop
- End

# Search Space

---

- For a simple function  $f(x)$ - Search space is one dimensional
- Encoding several values into the chromosome many dimensions can be searched  
e.g. two dimensions  $f(x, y)$
- Search space can be visualised as a surface or *fitness landscape* in which fitness dictates height
- Each possible genotype is a point in the space
- A GA tries to move the points to better places (*higher fitness*) in the space

# Search Space

---

- Nature of the search space dictates how a GA will perform
- A completely random space would be bad for a GA
- GA's can get stuck in local maxima if search spaces contain lots of these random spaces
- Generally, spaces in which *small improvements* get closer to the *global optimum* are good



# Example population


---

No.	Chromosome	Fitness
1	1010011010	1
2	1111100001	2
3	1011001100	3
4	1010000000	1
5	0000010000	3
6	1001011111	5
7	0101010101	1
8	1011100111	2

# GA operators: Selection

---

- Main idea: better individuals get higher chance
  - Chances proportional to fitness
  - Implementation: roulette wheel technique
    - Assign to each individual a part of the roulette wheel
    - Spin the wheel  $n$  times to select  $n$  individuals



fitness(A) = 3

fitness(B) = 1

fitness(C) = 2

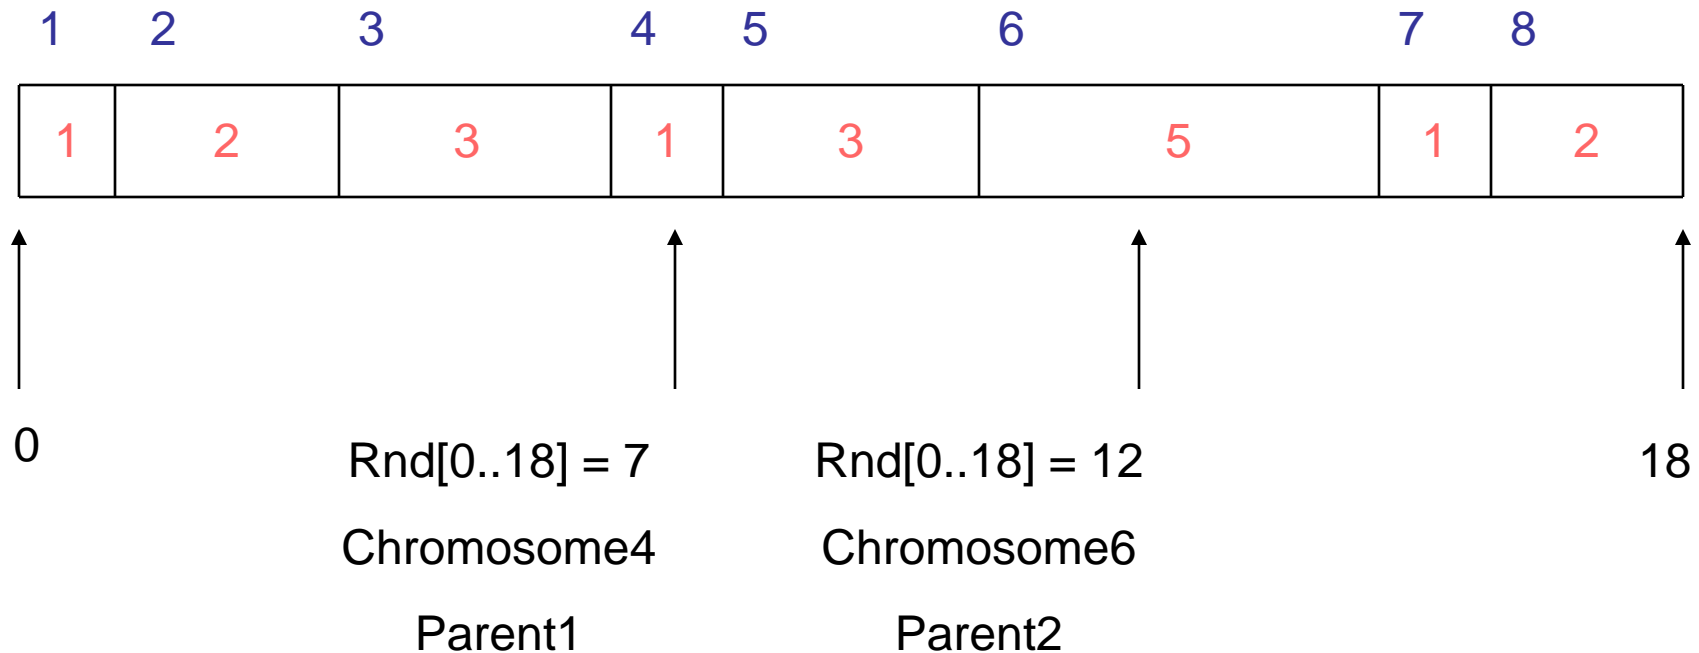
# GA operator: Selection

---

- Add up the fitness's of all chromosomes
- Generate a random number  $R$  in that range
- Select the first chromosome in the population that -when all previous fitness's are added - gives you at least the value  $R$

# Roulette Wheel Selection

---




# GA operator: Crossover

---

- Choose a random point on the two parents
- Split parents at this crossover point
- With some high probability (*crossover rate*) apply crossover to the parents
  - $P_c$  typically in range (0.6, 0.9)
- Create children by exchanging tails

# Crossover - Recombination

---




Crossover  
single point -  
random

*Single Point Crossover*

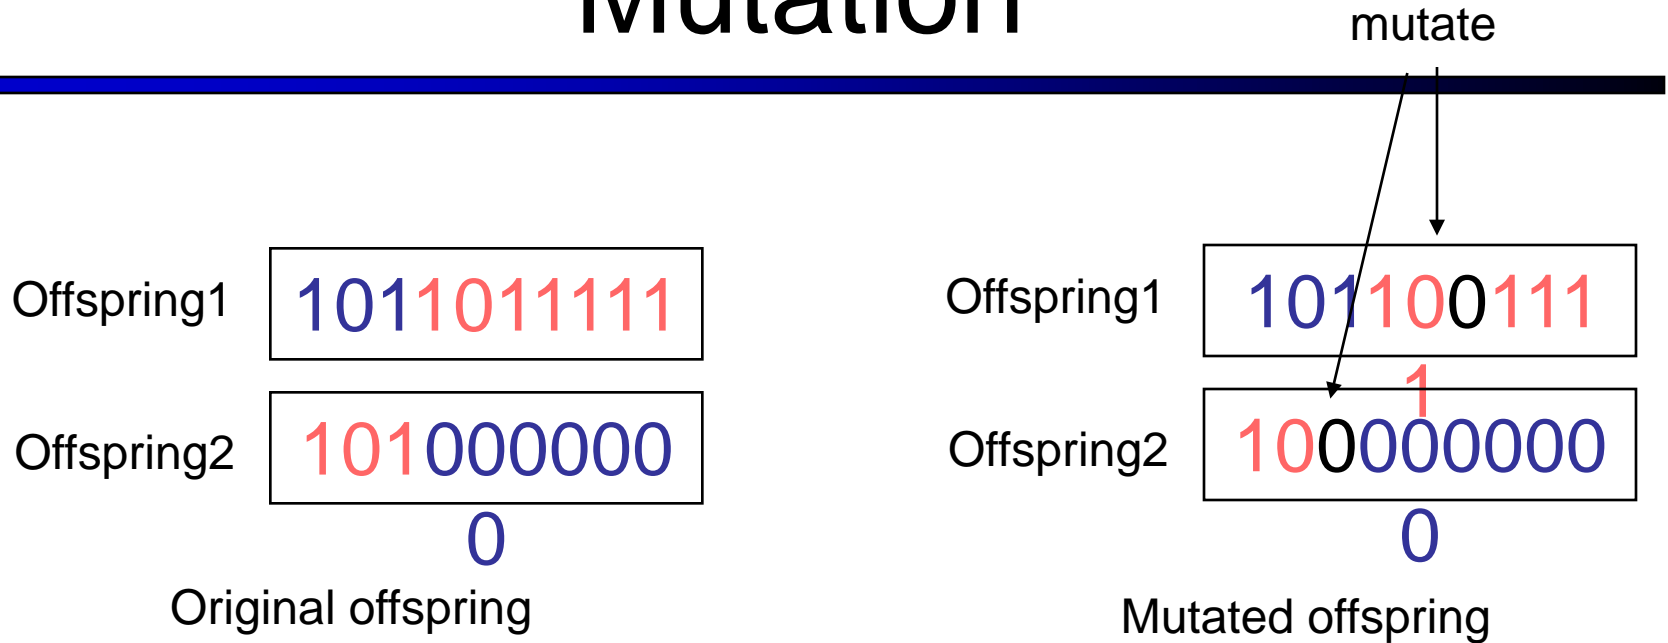
# n-point crossover

---

- Choose n random crossover points
- Split along those points
- Glue parts, alternating between parents
- Generalisation of 1 point (still some positional bias)



# Mutation



With some small probability (the *mutation rate*)  
flip each bit in the offspring (*typical values  
between 0.1 and 0.001*)



# Some points on GA

---

- Two quite different parents could generate a child that is long away from any of the parents
- More diversity of population at earlier stages
  - Similar to SA
  - Takes large steps early in the search process
  - Smaller steps later on when individuals are quite similar
- Combines uphill tendency with random exploration and exchanges information like stochastic beam search
  - Advantage of GA over stochastic beam search is due to crossover operation
- If random mutation at early stage: no advantage for crossover operation

# Some points on GA

---

- Advantage of crossover comes from the ability of crossover between large blocks
  - Schema
  - E.g. 246 constitutes a useful block
- Schema
  - 246\*\*\*\*\*
  - Instances: 24613578 ....
- Number of instances in population grows
  - If average fitness of all instances is greater than the mean
- Effect of schema is insignificant
  - Adjacent bits are completely different

*Thus, good chromosome representation is very important*

# GA: Quick Overview

---

- Typically applied to:
  - discrete optimization
- Attributed features:
  - not too fast
  - good heuristic for combinatorial problems
- Special Features:
  - Traditionally emphasizes combining information from good parents (*crossover*)
  - many variants, e.g., reproduction models, operators